

A WORKBOOK FROM THE SUPPLY CHAIN OF INTELLIGENCE™

The AI Defensibility Audit

Eight questions that separate a fortress from a wrapper. Score your product, your portfolio, or the company you're about to invest in — in under thirty minutes.

By Anand Arivukkarasu · Ex-Meta (Instagram) Product Leader · supplychainofai.com

How to use this workbook

Most AI strategy work confuses *traction* with *defensibility*. A product can be growing fast and still be one platform release away from extinction. The Defensibility Audit is a structural test — not a product test, not a market test. It asks the only question that matters in a world where the model itself is rented: **what about you cannot be absorbed?**

The mechanics

Score each of the eight questions **1 to 5**. Be honest — score what is true today, not what is on the roadmap. Sum to a total out of 40, then read the band on page 11. The bands are not grades. They are *moves*: each one tells you what to do next.

Three ways to run it

- **Solo, on your own product.** Thirty minutes. Brutal honesty required.
- **With your leadership team.** Score independently, then reconcile. The arguments are the value.
- **On a portfolio.** Run it on every AI company you own or are tracking. The spread is the alpha.

What this audit refuses to measure

Revenue. Growth rate. NPS. Fundraising. Team pedigree. None of these tell you whether the company will exist in three years when the platforms ship the obvious feature. That is the only question here.

The Supply Chain of Intelligence™ is a 10-layer model of the generative AI stack. This audit derives from it. The full framework, 19 worked case studies, and three structural-law essays are free at supplychainofai.com.

The ten layers, in one page

Every question in this audit maps to a layer (or two) of the Supply Chain of Intelligence. If you have not seen the framework before, this is the entire mental model:

Layer	Name	What it is
L-1	Land & Power	Energy, real estate, water — the substrate compute runs on.
L0	Compute & Silicon	GPUs, TPUs, custom silicon, datacenters.
L1	Data	The raw and proprietary corpora that models learn from.
L2	Models	Foundation models — the thing everyone now rents.
L3	Gates	Trust, safety, evals, compliance, verification.
L4	Distribution	Channels, app stores, OS integrations, enterprise relationships.
L5	Workflows	Encoded process — playbooks, agents, domain logic.
L6	Orchestration	Multi-step coordination, routing, planning.
L7	Surface	The UI. What the user actually touches.
L8	Memory	State, history, identity — the layer that compounds.

The asymmetry: most AI products live entirely at L7 (Surface). Defensibility comes from owning structural layers underneath — L1 (proprietary data), L3 (trust gates), L4 (distribution), L5 (workflow), L8 (memory). The audit is built to find out how many you actually own.

The eight questions

Each question has a 1 / 3 / 5 rubric. Scores of 2 and 4 are allowed — use them when you are between bands. The *common trap* under each question is the lie product leaders most often tell themselves about that layer. Read it before you score.

Q1

Model dependency

Layer: L2

Could a better GPT/Claude/Gemini release replace your core value?

Score 1 EXPOSED	Our core value is the model. If GPT-5 ships tomorrow with our feature, we are over.
Score 3 MIXED	Model improvements help us, but we add a layer on top — light playbooks, branding, a workflow.
Score 5 FORTIFIED	A frontier model release strengthens us. Our value lives in layers the model cannot reach.

Common trap: Confusing 'we use the best model' with 'we are protected from the best model.'

Your score

12345

Notes:

Q2

Data ownership

Layer: L1b

Do you create or own proprietary context that competitors can't access?

Score 1 EXPOSED	We use public data, prompts, or whatever the user types in. Anyone can recreate our corpus.
Score 3 MIXED	We have some unique data — a curated dataset, a labeled corpus, customer feedback — but it's replicable with effort.
Score 5 FORTIFIED	We generate data that exists nowhere else: tenant-private conversations, transaction history, sensor readings, behavioral graphs that compound with every user.

Common trap: Mistaking 'data we touch' for 'data we own.' If the user could export it and walk, you don't own it.

Your score

12345

Notes:

Q3

Workflow depth

Layer: L5 / L6

Are you embedded in a daily or high-stakes workflow users can't easily exit?

Score 1 EXPOSED	We're a tab. Users open us when they remember, and close us when they don't.
Score 3 MIXED	We're used weekly inside a real process, but the process would survive without us.
Score 5 FORTIFIED	We sit inside a daily, high-stakes, multi-step workflow. Removing us breaks the operating model.

Common trap: Counting MAU instead of counting load-bearing usage. A product can have high MAU and zero workflow depth.

Your score 1 2 3 4 5 _____ Notes: _____

Q4

Trust gate

Layer: L3

Do users rely on you for verification, compliance, quality, or approval?

Score 1 EXPOSED	We help users do a thing. They double-check the output themselves — or don't.
Score 3 MIXED	We have some QA / eval layer, but users could verify elsewhere.
Score 5 FORTIFIED	We are the verification. Regulators, auditors, or the user's customers accept our stamp. Replacing us means redoing compliance.

Common trap: Calling a confidence score a 'gate.' Gates have institutional acceptance behind them, not just a number.

Your score 1 2 3 4 5 _____ Notes: _____

Q5

Distribution

Layer: L4 / L7c

Do you own a channel, community, brand, or enterprise relationship?

Score 1 EXPOSED	We acquire users one paid click at a time. The channel belongs to Google or Meta.
Score 3 MIXED	We have a meaningful brand or community, but most growth still comes through rented channels.
Score 5 FORTIFIED	We own the railroad — preinstalled distribution, an enterprise contract, a developer ecosystem, an editorial brand competitors can't replicate.

Common trap: Confusing virality with distribution. Virality is a moment. Distribution is a structure.

Your score 1 2 3 4 5 _____ Notes: _____

Q6

Memory

Layer: L8

Does the product become smarter or more useful with usage history?

Score 1 EXPOSED	Every session starts from zero. The thousandth interaction is identical to the first.
Score 3 MIXED	We remember preferences, recent activity, light context — but the product is roughly the same at month one and month twelve.
Score 5 FORTIFIED	The product compounds. Month twelve is materially better than month one because of accumulated state — and that state cannot be ported elsewhere.

Common trap: Storing logs is not memory. Memory is state the product *uses* to get better.

Your score 1 2 3 4 5 _____ Notes: _____

Q7

Switching cost

Layer: L8d

Would leaving you destroy useful state, process, or institutional knowledge?

Score 1 EXPOSED	A user could switch over the weekend and nobody would notice on Monday.
Score 3 MIXED	Switching is annoying — some retraining, some lost configs — but doable in a quarter.
Score 5 FORTIFIED	Leaving means destroying months or years of compounded context, custom playbooks, integrations, and team-level fluency. The exit cost dwarfs the entry cost.

Common trap: Switching cost from *contract length* doesn't count. It evaporates at renewal.

Your score 1 2 3 4 5 _____ Notes: _____

Q8

Platform exposure

Layer: L2 / L7

Could a major platform (OpenAI, Google, Microsoft, Salesforce) bundle this for free?

Score 1 EXPOSED	Yes, and they probably will. We are the obvious next feature in someone's roadmap.
Score 3 MIXED	They could ship a thinner version, but enterprise customers would still want our depth.
Score 5 FORTIFIED	Even if a platform shipped a free version tomorrow, our customers would stay — because the layers we own (data, workflow, memory, trust) don't exist inside the platform.

Common trap: Believing 'the platform won't bother with us' is a strategy. It is not.

Your score 1 2 3 4 5 _____ Notes: _____

Your score sheet

Transfer your scores. Sum the column. Read your band on the next page.

#	Area	Layer	Score (1–5)
Q1	Model dependency	L2	_____
Q2	Data ownership	L1b	_____
Q3	Workflow depth	L5 / L6	_____
Q4	Trust gate	L3	_____
Q5	Distribution	L4 / L7c	_____
Q6	Memory	L8	_____
Q7	Switching cost	L8d	_____
Q8	Platform exposure	L2 / L7	_____
TOTAL			_____ / 40

Read your verdict

The band is not a grade — it's a *next move*. Find your total. Read what to do.

8–16

Thin Wrapper
The platform will absorb you.

What to do: Stop building features. Pick one layer below L7 — data, workflow, memory — and rebuild around it. If you cannot, sell now while the brand still has value.

Useful Tool, Weak Moat

Real utility, no structural protection.

17–24

What to do: You have time, but not much. Pick the single layer where you can credibly go deep in 18 months and concentrate every roadmap dollar there. Generic AI features are now a tax, not a moat.

Workflow Product

Embedded, survivable — watch the platforms.

25–30

What to do: Defend the workflow. Add a second layer — memory or trust gate — before a platform ships your category. Your enterprise contract is your shield; renew aggressively.

Defensible AI System

Multiple layers owned. The chain is yours.

31–36

What to do: Stop apologizing for price. Lean into the layers you own in your positioning. Recruit ahead of growth — the talent war for L1/L5/L8 operators is the next bottleneck.

Intelligence Gate

Platform candidate. You are the bottleneck.

37–40

What to do: Behave like a platform. Open APIs, fund an ecosystem, write the standards. Your risk is no longer being absorbed — it is being regulated. Get ahead of that conversation.

Three worked examples

The audit applied to three real companies. Same questions. Three very different verdicts. The point is not the scores — it is how quickly the structural picture clarifies.

Sierra

Fortress · 33/40

Customer-care agent built on proprietary tenant data, deep playbooks, and per-customer compounding memory.

Question	Score	Note
Model dependency	3	Uses frontier models, but value is not in the model.
Data ownership	5	Tenant conversations live behind their walls.
Workflow depth	4	Embedded in support operations — daily, load-bearing.
Trust gate	3	Brand voice + escalation rules act as a quality gate.
Distribution	2	Direct enterprise sales — no platform railroad.
Memory	5	Every resolution updates the customer-specific model.
Switching cost	5	Leaving destroys months of tenant-tuned behavior.
Platform exposure	3	OpenAI could ship a thin version; depth keeps enterprise.

Verdict: Hard to displace even by L4 owners. Triangle complete.

Jasper

Wrapper-at-risk · 12/40

Marketing copy surface with prompt templates and brand voice presets — no layer below the UI.

Question	Score	Note
Model dependency	1	Core value = the model. ChatGPT shipped, value collapsed.
Data ownership	1	User prompts; nothing exclusive.
Workflow depth	2	Used in marketing teams, but replaceable.
Trust gate	1	No gate. Output goes straight to human review.
Distribution	1	Standalone web app. No editor footprint.
Memory	1	Per-doc state. No org-level loop.
Switching cost	2	Some saved templates. Days, not quarters.
Platform exposure	3	Already realized — 80% mark-down confirmed.

Verdict: Commoditized the moment L2 shipped chat.

Glean

Ambiguous → tilting fortress · 26/40

Enterprise search that crawls every tenant SaaS and accumulates query-graph memory — but Copilot owns the OS.

Question	Score	Note
Model dependency	3	Model-agnostic by design.
Data ownership	4	Tenant-wide connector graph; deeper than Copilot for non-MSFT stacks.
Workflow depth	3	Daily use, but not yet load-bearing.
Trust gate	3	Permissions + ranking = encoded enterprise know-how.
Distribution	2	Strong inside IT — but Copilot is bundled in 365.
Memory	4	Per-employee answer graph compounds with usage.
Switching cost	4	Reconnecting every SaaS connector is a quarter of work.
Platform exposure	3	Copilot is the live threat — outcome decides 24 months.

Verdict: Triangle nearly complete. The L4 fight against Copilot decides the next 24 months.

Why this is not JTBD

Jobs-to-be-Done finds **demand** — why a user will hire your product. The Defensibility Audit finds **defensibility** — why a platform won't fire you next quarter. Both matter. They are not substitutes.

Question	JTBD	Supply Chain of Intelligence
What does it answer?	Why will users hire this product?	Why won't a platform fire it next quarter?
What does it find?	Demand.	Defensibility.
Time horizon	Today's user need.	Tomorrow's structural position.
Failure mode it catches	Building something nobody wants.	Building something everyone can copy or absorb.
Audience	PMs, designers, researchers.	Founders, product leaders, investors, boards.
Output	Roadmap, features, positioning.	Layer ownership, moat strategy, exit/defend/deepen call.

Keep going

This workbook is a slice of a larger framework. If the audit was useful, the rest is free at supplychainofai.com:

- The full 10-layer, 50-sublayer breakdown of the generative AI stack.
- 19 worked case studies — Sierra, Harvey, Gamma, Glean, Cursor, Perplexity, and more.
- Three structural-law essays: the laws that decide who wins the next decade.
- The Intelligence Cube — a 10×10×10 visual of where companies sit in the stack.

License. Creative Commons Attribution 4.0. Use this workbook with your team, your portfolio, your students, your podcast. Remix it. Argue with it. Translate it. The only ask: credit *Anand Arivukkarasu, Supply Chain of Intelligence™* (supplychainofai.com).

About the author. Anand Arivukkarasu is an ex-Meta (Instagram) product leader and AI product architect based in San Francisco. He has held VP/Head of Product roles at Ideas2IT, Refersion, and GRIN, and was earlier Lead PM at Vungle and Pinsight Media. He angel-invests in and advises early-stage AI companies. LinkedIn: linkedin.com/in/anandarivu